

Synchronism and Events

AdC - 2025

Command Queues

- Command queues are used to submit work to a device
- Two main types of command queues
 - In Order Queue
 - Out of Order Queue
- Used to measure the performance of an application as a whole.
- This necessitates understanding of OpenCL synchronization techniques and events

Command Queues

```
cl_command_queue clCreateCommandQueue (cl_context context,  
                                     cl_device_id device,  
                                     cl_command_queue_properties properties,  
                                     cl_int *errcode_ret)
```

creates a command-queue on a specific device.


context must be a valid OpenCL context.

Command-Queue Properties	Description
CL_QUEUE_OUT_OF_ORDER_EXEC_MODE_ENABLE	Determines whether the commands queued in the command-queue are executed in-order or out-of-order. If set, the commands in the command-queue are executed out-of-order. Otherwise, commands are executed in-order. For a detailed description about CL_QUEUE_OUT_OF_ORDER_EXEC_MODE_ENABLE , refer to <i>section 5.11</i> .
CL_QUEUE_PROFILING_ENABLE	Enable or disable profiling of commands in the command-queue. If set, the profiling of commands is enabled. Otherwise profiling of commands is disabled.

[Documentación PyOpenCL](#)


In-Order Execution

- In an in-order command queue, each command executes after the previous one has finished
 - For the set of commands shown, the read from the device would start after the kernel call has finished
- Memory transactions have consistent view



Out-of-Order Execution

- In an out-of-order command queue, commands are executed as soon as possible, without any ordering guarantees
- All memory operations occur in single memory pool
- Out-of-order queues result in memory transactions that will overlap and clobber data without some form of synchronization
- The commands discussed in the previous slide could execute in any order on device



Host Synchronization

- Synchronization is required if we use an out-of-order command queue or multiple command queues
- Synchronization is restricted to within a context
- This is similar to the fact that it is not possible to share data between multiple contexts without explicit copying
- The proceeding discussion of synchronization is applicable to any OpenCL device (CPU, GPU, etc)

Command Queue Control


- Command queue synchronization methods work on a per-queue basis
- **Finish:** `clFinish(cl_commandqueue)`
 - Waits for all commands in the command queue to complete before proceeding (**host blocks on this call**)
- **Barrier:** `clEnqueueBarrier(cl_commandqueue)`
 - Enqueue a synchronization point that ensures all prior commands in a queue have completed before any further commands execute

Synchronization for clEnqueue Functions

- Functions like `clEnqueueReadBuffer` and `clEnqueueWriteBuffer` have a boolean parameter to determine if the function is blocking
 - This provides a blocking construct that can be invoked to block the host
- If blocking is **TRUE**, OpenCL enqueues the operation using the host pointer in the command-queue
 - Host pointer **can** be reused by the application after the enqueue call returns
- If blocking is **FALSE**, OpenCL will use the host pointer parameter to perform a non-blocking read/write **and returns immediately**
 - Host pointer **cannot** be reused safely by the application after the call returns
 - Event handle returned by `clEnqueue*` operations can be used to check if the non-blocking operation has completed


Kernel Space

- In Kernel Space: “Device Synchronization model”



Device Synchronism

Kernel Synchronization: Barriers



Device Synchronism

```
// Kernel
__kernel void simpleKernel(
    __global float *a,
    __global float *b,
    __local float *l )
{
    l[get_local_id(0)] = a[get_global_id(0)];
    barrier(CLK_LOCAL_MEM_FENCE);
    unsigned int otherAddress =
        (get_local_id(0) + 1) % get_local_size(0);
    b[get_local_id(0)] = l[get_local_id(0)] + l[otherAddress];
}
```

OpenCL Events

- **Explicit synchronization is required for**
 - Out-of-order command queues
 - Multiple command queues
- OpenCL events are data-types defined by the specification for storing timing information returned by the device

OpenCL Events

- Profiling of OpenCL programs using events has to be enabled explicitly when creating a command queue
 - `CL_QUEUE_PROFILING_ENABLE` flag must be set
 - `command_queue = clCreateCommandQueue(context, devices[deviceUsed], CL_QUEUE_PROFILING_ENABLE, &err);`
 - Keeping track of events may slow down execution
- A handle to store event information can be passed for all `clEnqueue*` commands
 - When commands such as `clEnqueueNDRangeKernel` and `clEnqueueReadBuffer` are invoked timing information is recorded at the passed address

OpenCL Events

```
cl_int clEnqueueNDRangeKernel (cl_command_queue command_queue,  
                               cl_kernel kernel,  
                               cl_uint work_dim,  
                               const size_t *global_work_offset,  
                               const size_t *global_work_size,  
                               const size_t *local_work_size,  
                               cl_uint num_events_in_wait_list,  
                               const cl_event *event_wait_list,  
                               cl_event *event)
```

OpenCL Events

```
cl_int clEnqueueReadBuffer(cl_command_queue command_queue,  
                           cl_mem buffer,  
                           cl_bool blocking_read,  
                           size_t offset,  
                           size_t size,  
                           void *ptr,  
                           cl_uint num_events_in_wait_list,  
                           const cl_event *event_wait_list,  
                           cl_event *event)
```

```
cl_int clEnqueueWriteBuffer(cl_command_queue command_queue,  
                           cl_mem buffer,  
                           cl_bool blocking_write,  
                           size_t offset,  
                           size_t size,  
                           const void *ptr,  
                           cl_uint num_events_in_wait_list,  
                           const cl_event *event_wait_list,  
                           cl_event *event)
```

Uses of OpenCL Events

- Using OpenCL Events we can:
 - time execution of `clEnqueue*` calls like kernel execution or explicit data transfers
 - use the events from OpenCL to schedule asynchronous data transfers between host and device
 - observe overhead and time consumed by a kernel in the command queue versus actually executing
- **Note:** OpenCL event handling can be done in a consistent manner on both CPU and GPU for AMD and NVIDIA's implementations

Capturing Event Information

```
cl_int clGetEventProfilingInfo (
    cl_event event,           //event object
    cl_profiling_info param_name, //Type of data of event
    size_t param_value_size,    //size of memory pointed to by param_value
    void * param_value,        //Pointer to returned timestamp
    size_t * param_value_size_ret) //size of data copied to param_value
```

- `clGetEventProfilingInfo` allows us to query `cl_event` to get required counter values
- Timing information returned as `cl_ulong` data types
 - Returns device time counter in nanoseconds

Event Profiling Information

```
cl_int clGetEventProfilingInfo (
    cl_event event,           //event object
    cl_profiling_info param_name, //Type of data of event
    size_t param_value_size,   //size of memory pointed to by param_value
    void * param_value,       //Pointer to returned timestamp
    size_t * param_value_size_ret) //size of data copied to param_value
```

- Table shows event types described using `cl_profiling_info` enumerated type

Event Type	Description
CL_PROFILING_COMMAND_QUEUED	Command is enqueued in a command-queue by the host.
CL_PROFILING_COMMAND_SUBMIT	Command is submitted by the host to the device associated with the command queue.
CL_PROFILING_COMMAND_START	Command starts execution on device.
CL_PROFILING_COMMAND_END	Command has finished execution on device.

Profiling with Event Information

- Before getting timing information, we must make sure that the events we are interested in have completed
- There are different ways of waiting for events:
 - `clWaitForEvents(numEvents, eventlist)`
 - `clFinish(commandQueue)`
- Timer resolution can be obtained from the flag
`CL_DEVICE_PROFILING_TIMER_RESOLUTION`
when calling `clGetDeviceInfo()`

Event Profiling Information

```
command_queue = clCreateCommandQueue(context, devices[deviceUsed], CL_QUEUE_PROFILING_ENABLE, &err);

//Ensure to have executed all enqueued tasks
clFinish(command_queue);

//Launch Kernel linked to an event
err = clEnqueueNDRangeKernel(command_queue, kernel, 1, NULL, workGroupSize, NULL, 0, NULL, &event);

//Ensure kernel execution is finished
clWaitForEvents(1 , &event);

//Get the Profiling data
cl_ulong time_start, time_end;
double total_time;

clGetEventProfilingInfo(event, CL_PROFILING_COMMAND_START, sizeof(time_start), &time_start, NULL);
clGetEventProfilingInfo(event, CL_PROFILING_COMMAND_END, sizeof(time_end), &time_end, NULL);
total_time = time_end - time_start;
printf("\nExecution time in milliseconds = %0.3f ms\n", (total_time / 1000000.0) );
```

Using Events for Timing

- OpenCL events can easily be used for timing durations of kernels.
- This method is reliable for performance optimizations since it uses counters from the device
- By taking differences of the start and end timestamps we are discounting overheads like time spent in the command queue


```
clGetEventProfilingInfo( event_time,  
CL_PROFILING_COMMAND_START,  
sizeof(cl_ulong), &starttime, NULL);
```

```
clGetEventProfilingInfo(event_time,  
CL_PROFILING_COMMAND_END,  
sizeof(cl_ulong), &endtime, NULL);
```

```
unsigned long elapsed =  
(unsigned long)(endtime - starttime);
```

Ejercicio 3

Escribir un programa que realice dos veces la suma de los componentes vecinos de un arreglo 'A' como se muestra en la siguiente figura:



Utilizar los dos tipos de sincronismo: al nivel del host y al nivel del kernel.